

Stochastic Processes



Week 10 (Version 1.0)

Sampling Methods


Hamid R. Rabiee

Fall 2023

Overview

- Random Sampling
- Monte Carlo Principle
- Monte Carlo Markov Chain
- Metropolis Hasting
- Gibbs Sampling
- Monte Carlo & Nonparametric Bayesian models

Random Sampling



- Ω - “very large” sample set.
- π - probability distribution over Ω .

Goal: Sample points $x \in \Omega$ at random from distribution π .

The Probability Distribution


Typically,

$w:\Omega\rightarrow\mathbb{R}^+$ is an easily-computed weight function

$$\pi(x) = \frac{w(x)}{Z}$$

$Z = \sum_x w(x)$ is an unknown normalization factor

Example: Permutation of N Distinct Object



- Ω - all $N!$ permutations of N distinct objects.
- π - uniform distribution $[\forall \mathbf{x} \ w(\mathbf{x})=1]$.

Goal: pick a permutation uniformly at random.

Why Sampling?

- The use of samples allows us to conduct studies with more manageable data and in a timely manner.
- Randomly drawn samples do not have much bias if they are large enough, but achieving such a sample may be expensive and time-consuming.
- We often need to compute statistics of “typical” configurations: estimating mean of a stochastic process or mean energy, ...
- Estimating the statistics of a posterior density function in Bayesian inference.

Example: Estimating the Mean of $f(X)$

- Want to compute $E([f(X)])$ for function $f(\cdot)$.
- Standard method for approximating $E([f(X)])$ is to generate many **independent** sample values of X and compute sample mean of $f(X)$.
- Only useful in “trivial” cases where X can be generated directly.
- Many practical problems have non-trivial distribution for X
 - E.g., state in nonlinear/non-Gaussian state-space model, Bayesian inference, ...

The Monte Carlo principle

- **p(x)**: a target density defined over a high-dimensional space (e.g. the space of all possible configurations of a system under study)
- The idea of Monte Carlo techniques is to draw a set of (iid) samples $\{x^{(i)}\}$ for $i = 1, \dots, N$, from $p(x)$ in order to approximate $p(x)$ with the empirical distribution:

$$p(x) \approx \frac{1}{N} \sum_{i=1}^N \delta(x = x^{(i)})$$

- Using these samples we can approximate integrals $I(f)$ with tractable sums that converge (as the number of samples grows) to $I(f)$:

$$I(f) = \int f(x) p(x) dx \approx \frac{1}{N} \sum_{i=1}^N f(x^{(i)}) \xrightarrow{N \rightarrow \infty} I(f)$$

Importance sampling

- Target density $p(x)$ is known up to a constant
- Task: compute $I(f) = \int f(x)p(x)dx$

Idea:

- Introduce an arbitrary **proposal density** that includes the support of $p(x)$. Then:

$$I(f) = \int f(x) \underbrace{p(x)/q(x)}_{w(x) \text{ 'importance weight'}} * q(x)dx \approx \sum_{i=1}^N f(x^{(i)})w(x^{(i)})$$

- Sample from q instead of p
- Weight the samples according to their ‘importance’
- It also implies that $p(x)$ is approximated by:

$$p(x) \approx \sum_{i=1}^N w(x^{(i)})\delta(x = x^{(i)})$$

Efficiency depends on a ‘good’ choice of q .

Sequential Monte Carlo

- Sequential:
 - Real time processing
 - Dealing with non-stationarity
 - Not having to store the data
- Goal: estimate the distribution of ‘hidden’ trajectories:
 - We observe y_t at each time t : $p(x_{0:t} | y_{1:t})$, where
 - We have a model:
 - Initial distribution: $p(x_0)$
 - Dynamic model: $p(x_t | x_{0:t-1}, y_{1:t-1})$ for $t \geq 1$
 - Measurement model: $p(y_t | x_{0:t}, y_{1:t-1})$ for $t \geq 1$

Sequential Monte Carlo

- Can define a proposal distribution:

$$q(\tilde{x}_{0:t}|y_{1:t}) = p(x_{0:t-1}|y_{1:t-1})q(\tilde{x}_t|x_{0:t-1}, y_{1:t})$$

- Then the importance weights are:

$$\begin{aligned} w_t &= \frac{p(\tilde{x}_{0:t}|y_{1:t})}{q(\tilde{x}_{0:t}|y_{1:t})} = \frac{p(x_{0:t-1}|y_{1:t-1})}{p(x_{0:t-1}|y_{1:t-1})} \frac{p(\tilde{x}_t|x_{0:t-1}, y_{1:t})}{q(\tilde{x}_t|x_{0:t-1}, y_{1:t})} \\ &\propto \frac{p(y_t|\tilde{x}_t) p(\tilde{x}_t|x_{0:t-1}, y_{1:t-1})}{q_t(\tilde{x}_t|x_{0:t-1}, y_{1:t})}. \end{aligned}$$

- Simplifying the choice for proposal distribution:

Then: $q(\tilde{x}_t|x_{0:t-1}, y_{1:t}) = p(\tilde{x}_t|x_{0:t-1}, y_{1:t-1})$

$$w_t \propto p(y_t|\tilde{x}_t) \text{ ‘fitness’}$$

Sequential Monte Carlo

Sequential importance sampling step

- For $i = 1, \dots, N$, sample from the transition priors

$$\tilde{x}_t^{(i)} \sim q_t \left(\tilde{x}_t | x_{0:t-1}^{(i)}, y_{1:t} \right)$$

and set

$$\tilde{x}_{0:t}^{(i)} \triangleq \left(\tilde{x}_t^{(i)}, x_{0:t-1}^{(i)} \right)$$

- For $i = 1, \dots, N$, evaluate and normalize the importance weights


$$w_t^{(i)} \propto \frac{p \left(y_t | \tilde{x}_t^{(i)} \right) p \left(\tilde{x}_t^{(i)} | x_{0:t-1}^{(i)}, y_{1:t-1} \right)}{q_t \left(\tilde{x}_t^{(i)} | x_{0:t-1}^{(i)}, y_{1:t} \right)}.$$

Selection step

- Multiply/Discard particles $\left\{ \tilde{x}_{0:t}^{(i)} \right\}_{i=1}^N$ with high/low importance weights $w_t^{(i)}$ to obtain N particles $\left\{ x_{0:t}^{(i)} \right\}_{i=1}^N$.

Sequential Monte Carlo

$i=1, \dots, N=10$ particles



Three uses of Monte Carlo methods

1. For solving problems of probabilistic inference involved in developing computational models
2. As a source of hypotheses about how the mind might solve problems of probabilistic inference
3. As a way to explore people's subjective probability distributions


Applications on Monte Carlo Sampling

- Computer vision
- Speech & audio enhancement
- Web statistics estimation
- Regression & classification
- Bayesian networks
- Genetics & molecular biology
- Robotics, etc.
- ...

Markov chain Monte Carlo

- Basic idea: construct a *Markov chain* that will converge to the target distribution, and draw samples from that chain.
- Just uses something proportional to the target distribution (good for Bayesian inference!).
- Can work in state spaces of arbitrary (including unbounded) size (good for nonparametric Bayes).

Markov chains



Transition matrix

$$\mathbf{T} = P(\mathbf{x}^{(t+1)}|\mathbf{x}^{(t)})$$

Variables $\mathbf{x}^{(t+1)}$ independent of all previous variables given immediate predecessor $\mathbf{x}^{(t)}$

An example: card shuffling

- Each state $\mathbf{x}^{(t)}$ is a permutation of a deck of cards (there are $52!$ permutations)
- Transition matrix \mathbf{T} indicates how likely one permutation will become another
- The transition probabilities are determined by the shuffling procedure
 - riffle shuffle
 - overhand
 - one card

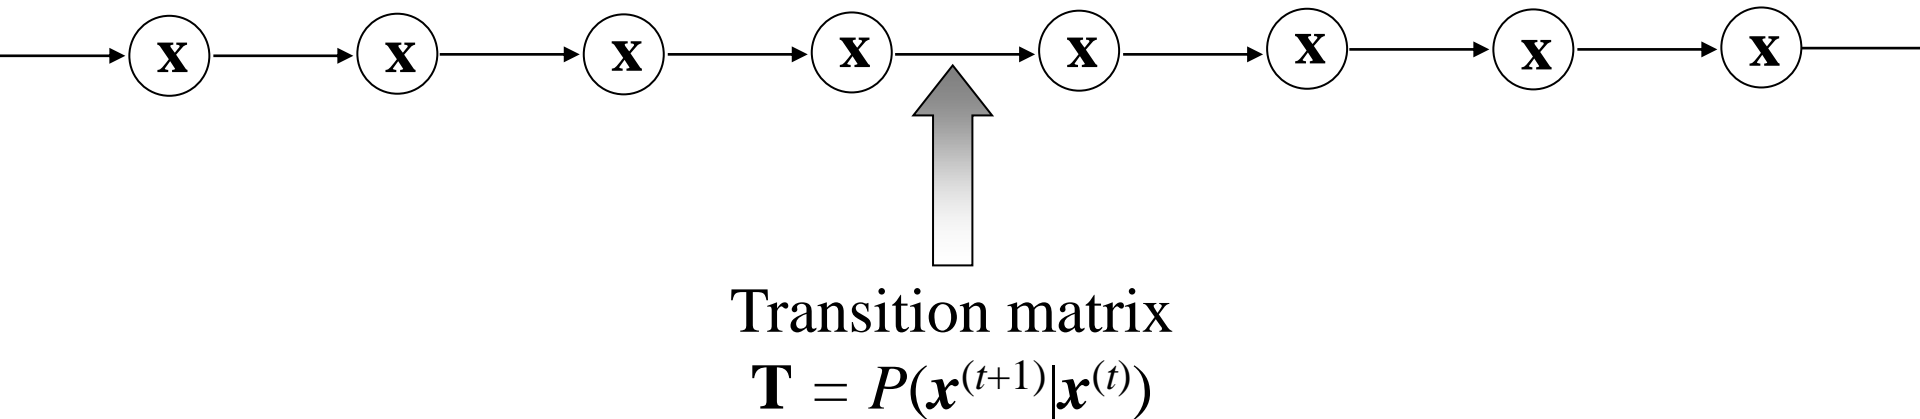
Convergence of Markov chains

- Why do we shuffle cards?
- Convergence to a uniform distribution takes only 7 riffle shuffles...
- Other Markov chains will also converge to a *stationary distribution*, if certain simple conditions are satisfied (called “ergodicity”)
 - e.g. every state can be reached in some number of steps from every other state

Modern Monte Carlo methods

- Sampling schemes for distributions with large state spaces known up to a multiplicative constant
- Two approaches:
 - importance sampling (and particle filters)
 - Markov chain Monte Carlo

Markov chain Monte Carlo



- States of chain are variables of interest
- Transition matrix chosen to give target distribution as stationary distribution

The Markov Chain Monte Carlo (MCMC)

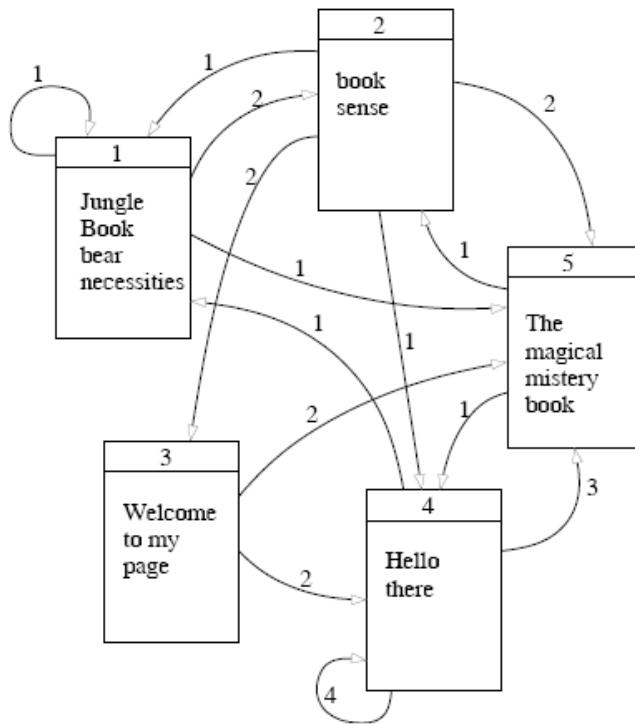
- Design a Markov Chain on finite state space:

state space: $x^{(i)} \in \{x_1, x_2, \dots, x_s\}$

Markov property: $p(x^{(i)} | x^{(i-1)}, \dots, x^{(1)}) = T(x^{(i)} | x^{(i-1)})$

such that when simulating a trajectory of states from it, it will explore the state space spending more time in the most important regions (i.e. where $p(x)$ is large)

Stationary distribution of a MC



- Suppose you browse this for infinitely long time, no matter where you started off:
- What is the probability to be at page x_i .
=>PageRank (Google)

$$p(x^{(i)} | x^{(i-1)}, \dots, x^{(1)}) = T(x^{(i)} | x^{(i-1)}) \equiv \mathbf{T}$$

$$((\mu(x^{(1)})\mathbf{T})\mathbf{T})\dots\mathbf{T} = \mu(x^{(1)})\mathbf{T}^n = p(x), \quad s.t. \quad p(x)\mathbf{T} = p(x)$$

Google vs. MCMC

$$p(x)\mathbf{T} = p(x)$$

- Google: given \mathbf{T} , finds $p(x)$
- MCMC: given $p(x)$, finds \mathbf{T}
 - But it also needs a ‘proposal (transition) probability distribution’ to be specified.
- Q: Do all MCs have a stationary distribution?
- A: No.

Conditions for existence of a unique stationary distribution

- Irreducibility
 - The transition graph is connected (any state can be reached)
- Aperiodicity
 - State trajectories drawn from the transition don't get trapped into cycles
- MCMC samplers are irreducible and aperiodic MCs that converge to the target distribution
- These 2 conditions are not easy to impose directly

Reversibility

- Reversibility (also called ‘detailed balance’) is a sufficient (but not necessary) condition for $p(x)$ to be the stationary distribution.

$$p(x^{(i)})T(x^{(i-1)}|x^{(i)}) = p(x^{(i-1)})T(x^{(i)}|x^{(i-1)}).$$

Summing both sides over $x^{(i-1)}$, gives us

$$p(x^{(i)}) = \sum_{x^{(i-1)}} p(x^{(i-1)})T(x^{(i)}|x^{(i-1)}).$$

- It is easier to work with this condition.

MCMC algorithms


- Metropolis-Hastings algorithm
- Metropolis algorithm
 - Mixtures and blocks
- Gibbs sampling
- other
- Sequential Monte Carlo & Particle Filters

Metropolis-Hastings algorithm


- Transitions have two parts:
 - proposal distribution: $q(\mathbf{x}^{(t+1)}|\mathbf{x}^{(t)})$
 - acceptance: take proposals with probability

$$A(\mathbf{x}^{(t)}, \mathbf{x}^{(t+1)}) = \min\left(1, \frac{P(\mathbf{x}^{(t+1)}) q(\mathbf{x}^{(t)}|\mathbf{x}^{(t+1)})}{P(\mathbf{x}^{(t)}) q(\mathbf{x}^{(t+1)}|\mathbf{x}^{(t)})}\right)$$


Metropolis-Hastings algorithm




Metropolis-Hastings algorithm




Metropolis-Hastings algorithm




Metropolis-Hastings algorithm



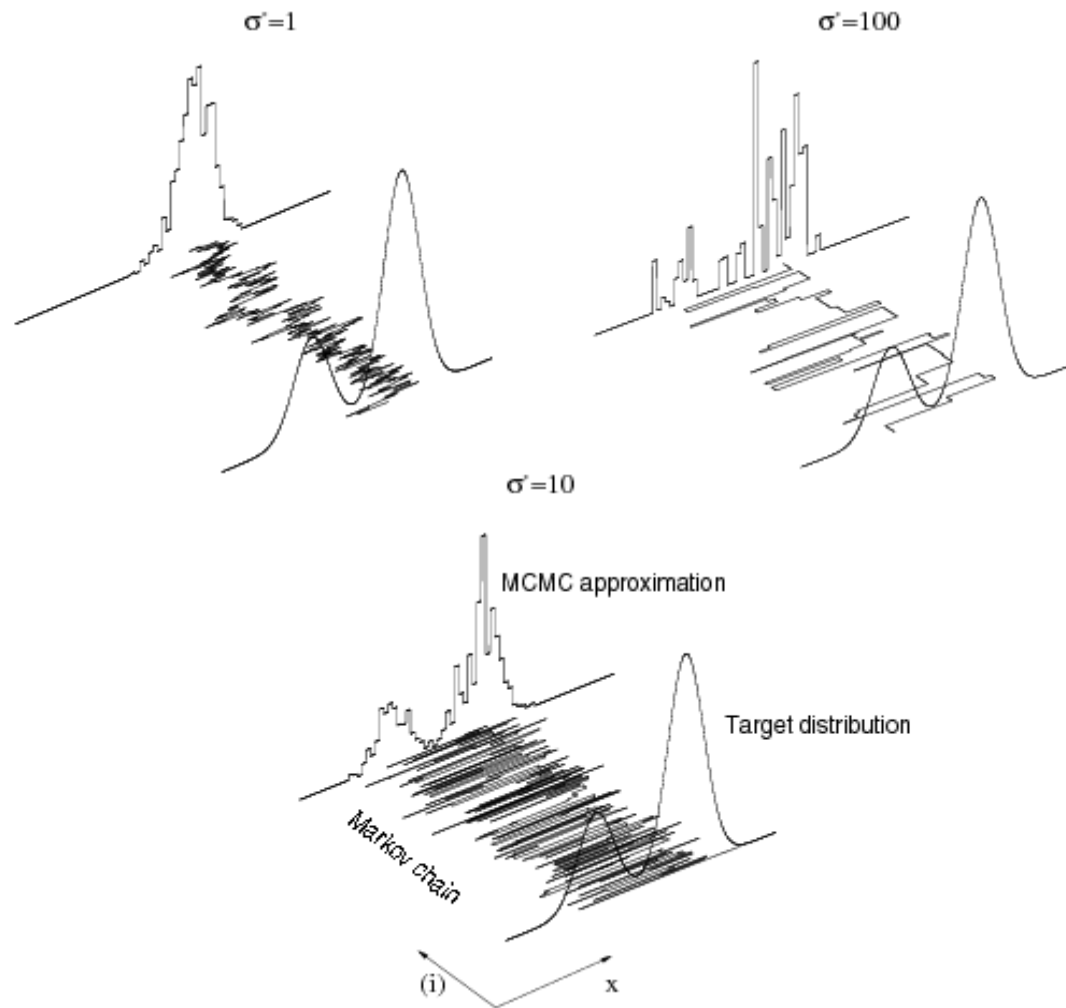
Metropolis-Hastings algorithm



Metropolis-Hastings algorithm



Examples of M-H simulations with q a Gaussian with variance σ



The Metropolis-Hastings and the Metropolis algorithm as a special case

1. Initialise $x^{(0)}$.
2. For $i = 0$ to $N - 1$
 - Sample $u \sim \mathcal{U}_{[0,1]}$.
 - Sample $x^* \sim q(x^*|x^{(i)})$.
 - If $u < \mathcal{A}(x^{(i)}, x^*) = \min\left\{1, \frac{p(x^*)q(x^{(i)}|x^*)}{p(x^{(i)})q(x^*|x^{(i)})}\right\}$
 $x^{(i+1)} = x^*$
else
 $x^{(i+1)} = x^{(i)}$

The Metropolis algorithm assumes a symmetric random walk proposal $q(x^*|x^{(i)}) = q(x^{(i)}|x^*)$ and, hence, the acceptance ratio simplifies to

$$\mathcal{A}(x^{(i)}, x^*) = \min\left\{1, \frac{p(x^*)}{p(x^{(i)})}\right\}.$$

Obs. The target distrib $p(x)$ is only needed up to normalisation.

Gibbs sampling

Gibbs sampling is a computationally convenient Bayesian inference algorithm that is a special case of the Metropolis–Hastings algorithm.

- Component-wise proposal q :

$$q(x^*|x^{(i)}) = \begin{cases} p(x_j^*|x_{-j}^{(i)}) & \text{If } x_{-j}^* = x_{-j}^{(i)} \\ 0 & \text{Otherwise.} \end{cases}$$

Where the notation means:

$$p(x_j|x_{-j}) = p(x_j|x_1, \dots, x_{j-1}, x_{j+1}, \dots, x_n)$$

- In this case, the acceptance probability is

$$\mathcal{A}(x^{(i)}, x^*) = 1$$

Gibbs Sampling

Particular choice of proposal distribution





For variables $\mathbf{x} = x_1, x_2, \dots, x_n$

Draw $x_i^{(t+1)}$ from $P(x_i/\mathbf{x}_{-i})$

$$\mathbf{x}_{-i} = x_1^{(t+1)}, x_2^{(t+1)}, \dots, x_{i-1}^{(t+1)}, x_{i+1}^{(t)}, \dots, x_n^{(t)}$$

(this is called the *full conditional* distribution)

Gibbs sampling



(MacKay, 2002)

Gibbs sampling algorithm

1. Initialise $x_{0,1:n}$.

2. For $i = 0$ to $N - 1$

– Sample $x_1^{(i+1)} \sim p(x_1 | x_2^{(i)}, x_3^{(i)}, \dots, x_n^{(i)})$.

– Sample $x_2^{(i+1)} \sim p(x_2 | x_1^{(i+1)}, x_3^{(i)}, \dots, x_n^{(i)})$.

\vdots

– Sample $x_j^{(i+1)} \sim p(x_j | x_1^{(i+1)}, \dots, x_{j-1}^{(i+1)}, x_{j+1}^{(i)}, \dots, x_n^{(i)})$.

\vdots

– Sample $x_n^{(i+1)} \sim p(x_n | x_1^{(i+1)}, x_2^{(i+1)}, \dots, x_{n-1}^{(i+1)})$.

The promise of particle filters

- People need to be able to update probability distributions over large hypothesis spaces as more data become available
- Particle filters provide a way to do this with limited computing resources:
 - maintain a fixed finite number of samples
- Not just for dynamic models:
 - can work with a fixed set of hypotheses, although this requires some further tricks for maintaining diversity

The magic of MCMC Methods

- Since we only ever need to evaluate the relative probabilities of two states, we can have huge state spaces (much of which we rarely reach)
- In fact, our state spaces can be *infinite*
 - common with nonparametric Bayesian models
- But... the guarantees it provides are asymptotic
 - making algorithms that converge in practical amounts of time is a significant challenge

References & Resources

- [1] M Isard & A Blake: CONDENSATION – conditional density propagation for visual tracking. J of Computer Vision, 1998.
- [2] C Andrieu, N de Freitas, A Doucet, M Jordan: An Introduction to MCMC for machine learning. Machine Learning, vol. 50, pp. 5--43, Jan. - Feb. 2003.
- [3] MCMC preprint service:
<http://www.statslab.cam.ac.uk/~mcmc/pages/links.html>.
- [4] W.R. Gilks, S. Richardson & D.J. Spiegelhalter: Markov Chain Monte Carlo in Practice. Chapman & Hall, 1996.
- Associated demos & further papers:
<http://www.robots.ox.ac.uk/~misard/condensation.html>.
- Nando de Freitas' MCMC papers & sw
<http://www.cs.ubc.ca/~nando/software.html>.

Next Weeks:

I hope you enjoyed this course!

Have a good Final Exam!