

# Stochastic Processes



**Week 10 (Version 1.0)**  
**Sampling Methods**


Hamid R. Rabiee

Fall 2025

# Overview

- Random Sampling
- Monte Carlo Principle
- Monte Carlo Markov Chain
- Metropolis Hasting
- Gibbs Sampling
- Monte Carlo & Nonparametric Bayesian models

# Random Sampling




- $\Omega$  - “very large” sample set.
- $\pi$  - probability distribution over  $\Omega$ .

Goal: Sample points  $x \in \Omega$  at random from distribution  $\pi$ .

# The Probability Distribution


Typically,

$w: \Omega \rightarrow \mathbb{R}^+$  is an easily-computed weight function

$$\pi(x) = \frac{w(x)}{Z}$$


$Z = \sum_x w(x)$  is an unknown normalization factor

# Example: Permutation of N Distinct Object



- $\Omega$  - all  $N!$  permutations of  $N$  distinct objects.
- $\pi$  - uniform distribution [ $\forall x w(x)=1$ ].

Goal: pick a permutation uniformly at random.

# Why Sampling?

- The use of samples allows us to conduct studies with more manageable data and in a timely manner.
- Randomly drawn samples do not have much bias if they are large enough, but achieving such a sample may be expensive and time-consuming.
- We often need to compute statistics of “typical” configurations: estimating mean of a stochastic process or mean energy, ...
- Estimating the statistics of a posterior density function in Bayesian inference.

# Example: Estimating the Mean of $f(X)$

- Want to compute  $E[f(X)]$  for function  $f(\cdot)$ .
- Standard method for approximating  $E[f(X)]$  is to generate many **independent** sample values of  $X$  and compute sample mean of  $f(X)$ .
- Only useful in “trivial” cases where  $X$  can be generated directly.
- Many practical problems have non-trivial distribution for  $X$ 
  - E.g., state in nonlinear/non-Gaussian state-space model, Bayesian inference, ...

# The Monte Carlo Principle

- $p(x)$ : a target density defined over a high-dimensional space (samples, the space of all possible configurations of a system under study)
- The idea of Monte Carlo techniques is to draw a set of (iid) samples  $\{x^{(i)}\}$  for  $i = 1, \dots, N$ , from  $p(x)$  in order to approximate  $p(x)$  with the empirical distribution:

$$p(x) \approx \frac{1}{N} \sum_{i=1}^N \delta(x = x^{(i)})$$

- Using these samples, we can approximate integrals  $I(f)$  with tractable sums that converge (as the number of samples grows) to  $I(f)$ :

$$I(f) = \int f(x)p(x)dx \approx \frac{1}{N} \sum_{i=1}^N f(x^{(i)}) \xrightarrow[N \rightarrow \infty]{} I(f)$$

# The Monte Carlo Core Concept

- The idea behind Monte Carlo sampling is to use randomness to approximate deterministic values.
- For example, If you want to estimate the value of a definite integral:

$$I = \int_a^b f(x)dx$$

- you can approximate this integral using Monte Carlo sampling by:

$$I \approx \frac{1}{N} \sum_{i=1}^N f(x_i)$$

- where  $x_i$  are random samples drawn uniformly from  $[a,b]$ .

# Steps & Benefits of Monte Carlo Sampling

## Steps in Monte Carlo Sampling:

- Define the Problem: Represent the problem in terms of probabilities or expectations.
- Generate Samples: Draw random samples from the relevant distribution or space.
- Compute Function Values: Evaluate the target function for each sample.
- Aggregate Results: Combine the results (e.g., averaging, summing) to estimate the desired quantity.

## Benefits of Monte Carlo Sampling

- **Scalability:** Works well in high-dimensional problems.
- **Flexibility:** Can handle complex, irregularly shaped domains or distributions.
- **Accuracy:** Improves with the number of samples, based on the law of large numbers.

# Example of Monte Carlo Sampling

## Example: Estimating $\pi$ Using Monte Carlo

- Imagine a circle inscribed in a square. The area of the circle ( $\pi r^2$ ) and the area of the square ( $4r^2$ ) give:

$$\frac{\text{Area of Circle}}{\text{Area of Square}} = \frac{\pi}{4}$$

To estimate  $\pi$ :

- Generate random points within the square.
- Count how many points fall inside the circle.
- Use the ratio of points in the circle to total points to estimate  $\pi$ :

$$\pi \approx 4 \times \frac{\text{Number of Points in Circle}}{\text{Total Number of Points}}$$

# Importance Sampling

- **Importance sampling** is a statistical technique for estimating properties of a particular distribution using samples from a different distribution.
- It is commonly used in situations where direct sampling from the target distribution is difficult or when some areas of the distribution contribute more significantly to the desired out


## Key Concepts

1. **Target Distribution ( $p(x)$ )**: The distribution you want to analyze or estimate properties for, such as mean or variance.
2. **Proposal Distribution ( $q(x)$ )**: The distribution from which samples are drawn because it is easier to sample from than the target distribution.

# Importance Sampling

## Core Idea

- Instead of sampling directly from  $p(x)$ , you sample from  $q(x)$  and reweight the samples to reflect their importance relative to  $p(x)$ . The weights correct for the fact that  $q(x)$  is not the target distribution.



# Importance Sampling

- **Estimating an Expectation**
- Suppose you want to estimate the expectation of a function  $f(x)$  under the target distribution  $p(x)$ :

$$\mathbb{E}_p[f(x)] = \int f(x)p(x)dx$$

- If sampling directly from  $p(x)$  is difficult, but you can sample from  $q(x)$ , the expectation can be rewritten as:

$$\mathbb{E}_p[f(x)] = \int f(x)\frac{p(x)}{q(x)}q(x)dx$$

- This implies:

$$\mathbb{E}_p[f(x)] \approx \frac{1}{N} \sum_{i=1}^N f(x_i)w(x_i)$$

# Importance Sampling

where:

- $x_i \sim q(x)$
- $w(x_i) = \frac{p(x_i)}{q(x_i)}$  are the importance weights.
- Normalized Weights to ensure stability and better convergence:

$$\mathbb{E}_p[f(x)] \approx \frac{\sum_{i=1}^N f(x_i)w(x_i)}{\sum_{i=1}^N w(x_i)}$$

Applications

- Monte Carlo Integration: Estimate integrals where direct evaluation is difficult.
- Bayesian Inference: Approximating posterior distributions in cases with complex likelihoods.
- Reinforcement Learning: Adjusting for the difference between behavior and target policies.

# Importance Sampling

## Challenges

- **Choice of Proposal Distribution:**  $q(x)$  must adequately cover the regions where  $p(x)$  is significant. Poor choices can lead to large variances in the estimates.
- **Weight Calculation:** Large differences between  $p(x)$  and  $q(x)$  can result in extreme weights, leading to numerical instability.

## Tips

- Ensure  $q(x)$  has heavier tails than  $p(x)$  to avoid zero weights.
- Normalize weights to improve stability.

# Sequential Monte Carlo

- Advantages of Sequential Sampling:
  - Real time processing
  - Dealing with non-stationarity
  - Not having to store the data
- Goal: Estimate the distribution of ‘hidden’ trajectories:
  - We observe  $y_t$  at each time  $t$ :  $p(x_{0:t} | y_{1:t})$ , where
  - We have a model:
    - Initial distribution:  $p(x_0)$
    - Dynamic model:  $p(x_t | x_{0:t-1}, y_{1:t-1})$  for  $t \geq 1$
    - Measurement model:  $p(y_t | x_{0:t}, y_{1:t-1})$  for  $t \geq 1$

# Sequential Monte Carlo

- Can define a proposal distribution:

$$q(\tilde{x}_{0:t}|y_{1:t}) = p(x_{0:t-1}|y_{1:t-1})q(\tilde{x}_t|x_{0:t-1}, y_{1:t})$$

- Then the importance weights are:

$$\begin{aligned} w_t &= \frac{p(\tilde{x}_{0:t}|y_{1:t})}{q(\tilde{x}_{0:t}|y_{1:t})} = \frac{p(x_{0:t-1}|y_{1:t})}{p(x_{0:t-1}|y_{1:t-1})} \frac{p(\tilde{x}_t|x_{0:t-1}, y_{1:t})}{q(\tilde{x}_t|x_{0:t-1}, y_{1:t})} \\ &\propto \frac{p(y_t|\tilde{x}_t) p(\tilde{x}_t|x_{0:t-1}, y_{1:t-1})}{q_t(\tilde{x}_t|x_{0:t-1}, y_{1:t})}. \end{aligned}$$

- Simplifying the choice for proposal distribution:

Then:  $q(\tilde{x}_t|x_{0:t-1}, y_{1:t}) = p(\tilde{x}_t|x_{0:t-1}, y_{1:t-1})$

$w_t \propto p(y_t|\tilde{x}_t)$  ‘fitness’

# Sequential Monte Carlo

## Sequential importance sampling step

- For  $i = 1, \dots, N$ , sample from the transition priors

$$\tilde{x}_t^{(i)} \sim q_t \left( \tilde{x}_t | x_{0:t-1}^{(i)}, y_{1:t} \right)$$

and set

$$\tilde{x}_{0:t}^{(i)} \triangleq \left( \tilde{x}_t^{(i)}, x_{0:t-1}^{(i)} \right)$$

- For  $i = 1, \dots, N$ , evaluate and normalize the importance weights


$$w_t^{(i)} \propto \frac{p \left( y_t | \tilde{x}_t^{(i)} \right) p \left( \tilde{x}_t^{(i)} | x_{0:t-1}^{(i)}, y_{1:t-1} \right)}{q_t \left( \tilde{x}_t^{(i)} | x_{0:t-1}^{(i)}, y_{1:t} \right)}.$$

## Selection step

- Multiply/Discard particles  $\left\{ \tilde{x}_{0:t}^{(i)} \right\}_{i=1}^N$  with high/low importance weights  $w_t^{(i)}$  to obtain  $N$  particles  $\left\{ x_{0:t}^{(i)} \right\}_{i=1}^N$ .

# Sequential Monte Carlo

$i=1, \dots, N=10$  particles



# Three uses of Monte Carlo methods

1. For solving problems of probabilistic inference involved in developing computational models
2. As a source of hypotheses about how the mind might solve problems of probabilistic inference
3. As a way to explore people's subjective probability distributions


# Applications of Monte Carlo Sampling

- Computer vision
- Speech & audio enhancement
- Web statistics estimation
- Regression & classification
- Bayesian networks
- Genetics & molecular biology
- Robotics, etc.
- ...

# Markov Chain Monte Carlo

- Basic idea: Construct a *Markov chain* that will converge to the target distribution and draw samples from that chain.
- Just uses something proportional to the target distribution (good for Bayesian inference!).
- Can work in state spaces of arbitrary (including unbounded) size (good for nonparametric Bayes).

# Markov Chains



Variables  $x^{(t+1)}$  independent of all previous variables given immediate predecessor  $x^{(t)}$

# An example: Card Shuffling

- Each state  $\mathbf{x}^{(t)}$  is a permutation of a deck of cards (there are  $52!$  permutations)
- Transition matrix  $\mathbf{T}$  indicates how likely one permutation will become another
- The transition probabilities are determined by the shuffling procedure:
  - Riffle Shuffle
  - Overhand
  - One Card


# Convergence of Markov Chains

- Why do we shuffle cards?
- Convergence to a uniform distribution takes only 7 riffle shuffles...
- Other Markov chains will also converge to a *stationary distribution*, if certain simple conditions are satisfied (called “ergodicity”)
  - e.g. every state can be reached in some number of steps from every other state

# Modern Monte Carlo methods

- Sampling schemes for distributions with large state spaces known up to a multiplicative constant
- Two approaches:
  - Importance Sampling (particle filters)
  - Markov Chain Monte Carlo

# Markov chain Monte Carlo



- States of chain are variables of interest
- Transition matrix chosen to give target distribution as stationary distribution

# The Markov Chain Monte Carlo (MCMC)


- Design a Markov Chain on finite state space:

state space :  $x^{(i)} \in \{x_1, x_2, \dots, x_s\}$

Markov property :  $p(x^{(i)} | x^{(i-1)}, \dots, x^{(1)}) = T(x^{(i)} | x^{(i-1)})$

such that when simulating a trajectory of states from it, it will explore the state space spending more time in the most important regions (i.e. where  $p(x)$  is large)

# Stationary distribution of a MC



- Suppose you browse the Web for an infinitely long time, no matter where you started:
- What is the probability of being at page  $x_i$ .  
=>**PageRank (Google)**

$$p(x^{(i)} | x^{(i-1)}, \dots, x^{(1)}) = T(x^{(i)} | x^{(i-1)}) \equiv \mathbf{T}$$

$$((\mu(x^{(1)})\mathbf{T})\mathbf{T})\dots\mathbf{T} = \mu(x^{(1)})\mathbf{T}^n = p(x), \quad s.t. \quad p(x)\mathbf{T} = p(x)$$

# Google vs. MCMC

$$p(x)\mathbf{T} = p(x)$$

- Google: given  $\mathbf{T}$ , finds  $p(x)$
- MCMC: given  $p(x)$ , finds  $\mathbf{T}$ 
  - But it also needs a ‘proposal (transition) probability distribution’ to be specified.
- Q: Do all MCs have a stationary distribution?
- A: No.

# Conditions for existence of a unique stationary distribution

- Irreducibility
  - The transition graph is connected (any state can be reached)
- Aperiodicity
  - State trajectories drawn from the transition don't get trapped into cycles
- MCMC samplers are irreducible and aperiodic MCs that converge to the target distribution
- These 2 conditions are not easy to impose directly

# Reversibility

- Reversibility (also called ‘detailed balance’) is a sufficient (but not necessary) condition for  $p(x)$  to be the stationary distribution.

$$p(x^{(i)})T(x^{(i-1)}|x^{(i)}) = p(x^{(i-1)})T(x^{(i)}|x^{(i-1)}).$$

Summing both sides over  $x^{(i-1)}$ , gives us

$$p(x^{(i)}) = \sum_{x^{(i-1)}} p(x^{(i-1)})T(x^{(i)}|x^{(i-1)}).$$

- It is easier to work with this condition.

# MCMC Algorithms


- Metropolis-Hastings algorithm
- Metropolis algorithm
  - Mixtures and blocks
- Gibbs sampling
- Sequential Monte Carlo & Particle Filters

# Metropolis-Hastings Algorithm


- Transitions have two parts:
  - proposal distribution:  $q(\mathbf{x}^{(t+1)}|\mathbf{x}^{(t)})$
  - acceptance: take proposals with probability

$$A(\mathbf{x}^{(t)}, \mathbf{x}^{(t+1)}) = \min\left( 1, \frac{P(\mathbf{x}^{(t+1)}) q(\mathbf{x}^{(t)}|\mathbf{x}^{(t+1)})}{P(\mathbf{x}^{(t)}) q(\mathbf{x}^{(t+1)}|\mathbf{x}^{(t)})} \right)$$


# Metropolis-Hastings algorithm




# Metropolis-Hastings algorithm




# Metropolis-Hastings algorithm




# Metropolis-Hastings algorithm




# Metropolis-Hastings algorithm



# Metropolis-Hastings algorithm



# Examples of M-H simulations with $q(x)$ a Gaussian with variance $\sigma$



# The Metropolis-Hastings and the Metropolis algorithm as a special case

```
1. Initialise  $x^{(0)}$ .  
2. For  $i = 0$  to  $N - 1$   
    – Sample  $u \sim \mathcal{U}_{[0,1]}$ .  
    – Sample  $x^* \sim q(x^*|x^{(i)})$ .  
    – If  $u < \mathcal{A}(x^{(i)}, x^*) = \min\left\{1, \frac{p(x^*)q(x^{(i)}|x^*)}{p(x^{(i)})q(x^*|x^{(i)})}\right\}$   
         $x^{(i+1)} = x^*$   
    else  
         $x^{(i+1)} = x^{(i)}$ 
```

The Metropolis algorithm assumes a symmetric random walk proposal  $q(x^*|x^{(i)}) = q(x^{(i)}|x^*)$  and, hence, the acceptance ratio simplifies to


$$\mathcal{A}(x^{(i)}, x^*) = \min\left\{1, \frac{p(x^*)}{p(x^{(i)})}\right\}.$$

Obs. The target distrib  $p(x)$  is only needed up to normalization.

# The Metropolis-Hastings Tutorial Clip

**Markov Chain Monte Carlo**

**Arianna Rosenbluth**



The diagram illustrates the Metropolis-Hastings algorithm. It features a plot of the target distribution  $f(x)$  against  $x$ , which is a mixture of two normal distributions:  $X \sim 0.5 \cdot \mathcal{N}(-2, 1) + 0.5 \cdot \mathcal{N}(2, 1)$ . A blue curve represents the proposal distribution  $\mathcal{N}(x_t, 0.4)$ , centered at  $x_t = 0.6$ . A yellow arrow points from the current sample at  $x=1$  towards a proposed sample at  $x=2$ . A legend box titled "Move Criterion" contains the following steps:  
1.  $x_{t+1} \sim \mathcal{N}(x_t, 0.4)$   
2.  $\alpha = \frac{f(x_{t+1})}{f(x_t)}$   
3.  $u \sim \mathcal{U}(0, 1)$   
4. if  $\alpha > u$   
    move to proposed sample  
else  
    stay at current sample

Should we move to this sample?

# Gibbs Sampling

Gibbs sampling is a computationally convenient Bayesian inference algorithm that is a special case of the Metropolis–Hastings algorithm.

- Component-wise proposal  $q(x)$ :

$$q(x^*|x^{(i)}) = \begin{cases} p(x_j^*|x_{-j}^{(i)}) & \text{If } x_{-j}^* = x_{-j}^{(i)} \\ 0 & \text{Otherwise.} \end{cases}$$

Where the notation means:

$$p(x_j|x_{-j}) = p(x_j|x_1, \dots, x_{j-1}, x_{j+1}, \dots, x_n)$$

- In this case, the acceptance probability is

$$\mathcal{A}(x^{(i)}, x^*) = 1$$

# Gibbs Sampling

Particular choice of proposal distribution





For variables  $\mathbf{x} = x_1, x_2, \dots, x_n$

Draw  $x_i^{(t+1)}$  from  $P(x_i | \mathbf{x}_{-i})$

$\mathbf{x}_{-i} = x_1^{(t+1)}, x_2^{(t+1)}, \dots, x_{i-1}^{(t+1)}, x_{i+1}^{(t)}, \dots, x_n^{(t)}$

(this is called the *full conditional* distribution)

# Gibbs sampling



(MacKay, 2002)

# Gibbs sampling algorithm

1. Initialise  $x_{0,1:n}$ .
2. For  $i = 0$  to  $N - 1$ 
  - Sample  $x_1^{(i+1)} \sim p(x_1 | x_2^{(i)}, x_3^{(i)}, \dots, x_n^{(i)})$ .
  - Sample  $x_2^{(i+1)} \sim p(x_2 | x_1^{(i+1)}, x_3^{(i)}, \dots, x_n^{(i)})$ .
  - ⋮
  - Sample  $x_j^{(i+1)} \sim p(x_j | x_1^{(i+1)}, \dots, x_{j-1}^{(i+1)}, x_{j+1}^{(i)}, \dots, x_n^{(i)})$ .
  - ⋮
  - Sample  $x_n^{(i+1)} \sim p(x_n | x_1^{(i+1)}, x_2^{(i+1)}, \dots, x_{n-1}^{(i+1)})$ .

# The Promise of Particle Filters

- People need to be able to update probability distributions over large hypothesis spaces as more data becomes available
- Particle filters provide a way to do this with limited computing resources:
  - Maintain a fixed finite number of samples
- Not just for dynamic models:
  - Can work with a fixed set of hypotheses, although this requires some further tricks for maintaining diversity

# The Magic of MCMC Methods

- Since we only ever need to evaluate the relative probabilities of two states, we can have huge state spaces (much of which we rarely reach)
- In fact, our state spaces can be *infinite*
  - Common with nonparametric Bayesian models
- But... the guarantees it provides are asymptotic
  - Making algorithms that converge in practical amounts of time is a significant challenge

# References & Resources

- [1] M Isard & A Blake: CONDENSATION – conditional density propagation for visual tracking. J of Computer Vision, 1998.
- [2] C Andrieu, N de Freitas, A Doucet, M Jordan: An Introduction to MCMC for machine learning. Machine Learning, vol. 50, pp. 5--43, Jan. - Feb. 2003.
- [3] MCMC preprint service:  
<http://www.statslab.cam.ac.uk/~mcmc/pages/links.html>.
- [4] W.R. Gilks, S. Richardson & D.J. Spiegelhalter: Markov Chain Monte Carlo in Practice. Chapman & Hall, 1996.
- Associated demos & further papers:  
<http://www.robots.ox.ac.uk/~misard/condensation.html>.
- Nando de Freitas' MCMC papers & sw  
<http://www.cs.ubc.ca/~nando/software.html>.

**Next Weeks:**

**I hope you enjoyed this course!**

**Have a good Final Exam!**